

Convert this docx to PDF before submitting

Should you use granger1 or granger2?

<https://fxlin.github.io/p2-concurrency/#which-server>

For p2, this is a soft suggestion, not a hard requirement.

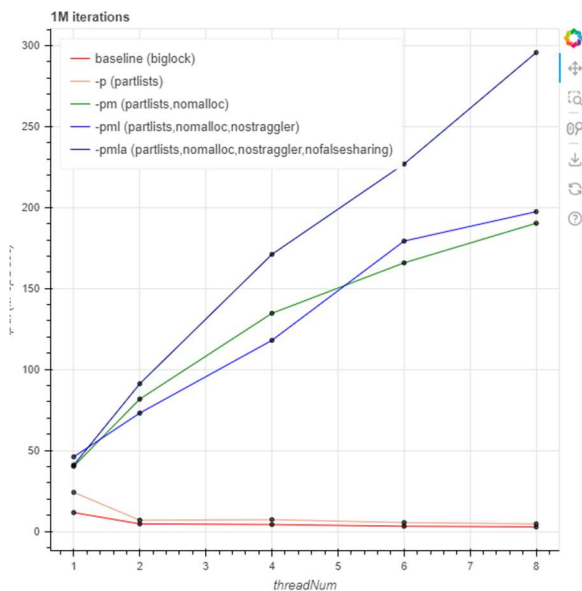
Scalability

0. Reproduce the benchmarks

Repeat what has been described in the project description.

- Attach a scalability plot (ONLY the one showing all the program versions) you generated. (10)

To generate the plot, you can use any tool. There's a boilerplate script (p2-concurrency/scripts/plot.py) that may help; but you are not required to use it.



- Compare your observation with the given results. What are the same? What are different? (10)
 - The observations are very similar to each other. The general shape of each line is very similar for both the generated graphs and the given results; that is, in both graphs, the -pm, -pml, and -pmla generally follow a positive linear slope and baseline and -p generally plateau. There are a couple of differences between

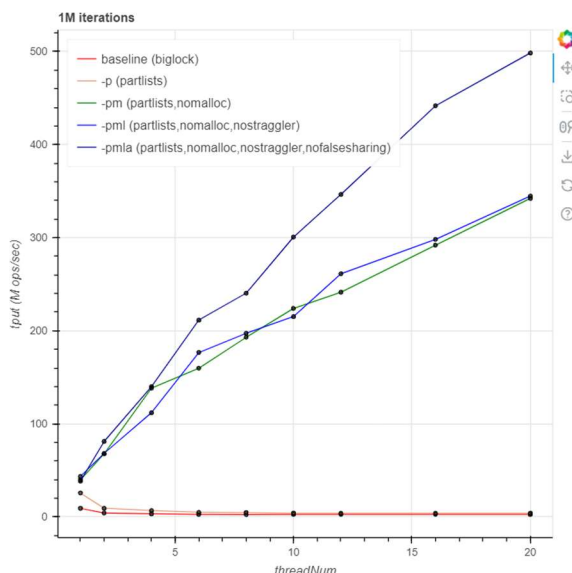
both graphs. The `-p` program seems to **drop less** between the **first two threadNum** in the generated graph than it does in the given graph and there seems to have **less noise** in the generated graph. The `-pm` programs seem to have a **higher throughput** than the `-pml` in the generated graph than the given graph. The `-pmla` program seems to generally have a **higher slope** in the generated graph than the given graph.

- Explain your observation. (10)
 - The graphs are similar because the **benefits** of implementing the improvements is **shared** regardless of the settings or configuration of the machines. In other words, we will always see a general linear positive trend for `-pmla` unless there is a **DRASTIC** change in the **configuration** and **settings** of the machine when the given results were created and when the generated results were created. At the same time, the difference between the two graphs are explained by the **simple differences** between the **configuration** and **settings** of the machine that generated the generated results versus the machine that generated the given results.

1. The unfinished scalability quest

How does the program scale to more than 8 cores?

- Attach a scalability plot (ONLY the one showing all the program versions) with core count = {1 2 4 6 8 10 12 16 20}. You may want to tweak `run.sh` and `plot.py` (10)



- Describe and explain your observation. (10)
 - It seems as if **baseline** and **-p** had an ok throughput for the **1-2 threads** and essentially plateaued around **0** at **higher threads**. The **baseline** and **-p** do not

scale well with higher threads because the throughput issue is more correlated with issues with the **expensive computation of `malloc()`**, processes waiting for **stragglers**, and **false sharing**. **Specializing memory allocation (`-pm`)** and assigning the **same amount of work to every worker (`-pml`)** both target a similar **facet** of the problem, that is, **allocation of tasks** to various worker threads. This explains why the **`-pm`** and **`-pml`** scale similarly with higher threads. Both of these programs scale better than **baseline** and **`-p`** because with higher threads more worker threads will **share and fight** for the **same resources** calling **`malloc()`** many times and creating a **bottleneck** for more **time-consuming tasks**. There is a **significant boost** in the **`-pmla`** scaling because the program addresses **another facet** of a lower throughput, that is, **cache misses** which can **drastically increase runtime**.

- If there's any scalability bottleneck, profile the execution with VTune (e.g. consider trying VTune's "microarchitecture exploration"). Can you make the program scale better? If so, show your code and profiling results; if not, reason about possible bottlenecks. (5)
 - A possible bottleneck could be a **specific function** in the program since the CPU Time seems to be present in a **specific portion** of the entire process as shown in the figure below from calling with parameter **`-iterations=1M -thread=8 -parts=32`**. Another bottleneck could be with the **data structure** that is used which is linear in insertion. The **workload** of individual threads could be made to be **more equal** with each other.

